

Bike sharing

Load default libraries

```
library(data.table)
library(dplyr)
library(tm)
library(ggplot2)
library(caret)
library(stringi)
library(FeatureHashing)
library(xgboost)
theme_set(theme_bw())
set.seed(42)
```

Read the data

```
train = fread("./prepared_train.csv", header = T, sep = ",", integer64 = "numeric")
validate = fread("./prepared_validate.csv", header = T, sep = ",", integer64 = "numeric")
test = fread("./prepared_test.csv", header = T, sep = ",", integer64 = "numeric")

test2 = fread("./test.csv", header = T, sep = ",", integer64 = "numeric")
```

Change type to categorical variables

```
fix_types = function(data) {
  data %>% mutate(
    month = as.character(month),
    day = as.character(day),
    weekday = as.character(weekday),
    hour = as.character(hour),
    season = as.character(season),
    holiday = as.character(holiday),
    workingday = as.character(workingday),
    weather = as.character(weather))
}
train = fix_types(train)
validate = fix_types(validate)
test = fix_types(test)
```

Prepare features

Prepare target

We going to predict $\log(Y+1)$ to optimize the target cost function

```

preparedTrainTarget = log(train$count + 1)
preparedValidateTarget = log(validate$count + 1)
preparedFullTarget = c(preparedTrainTarget, preparedValidateTarget)

```

Xgboost train and cross-validate

```

dtrain <- xgb.DMatrix(train_matrix, label = preparedTrainTarget)

model = xgboost(dtrain, nround=100, nthread = 2, nfold = 5, metrics=list("rmse"),
               max.depth = 3, eta = 0.1, objective = "reg:linear", print.every.n = 33)

```

```

## [0] train-rmse:3.910831
## [33] train-rmse:0.755147
## [66] train-rmse:0.603379
## [99] train-rmse:0.529682

```

```

history <- xgb.cv(data = dtrain, nround=150, nthread = 2, nfold = 5, metrics=list("rmse"),
                 max.depth = 3, eta = 0.1, objective = "reg:linear", print.every.n = 33)

```

```

## [0] train-rmse:3.910832+0.005531 test-rmse:3.910809+0.025410
## [33] train-rmse:0.750372+0.004117 test-rmse:0.762903+0.011960
## [66] train-rmse:0.603587+0.006684 test-rmse:0.621172+0.014271
## [99] train-rmse:0.532241+0.010550 test-rmse:0.552558+0.012204
## [132] train-rmse:0.492933+0.008041 test-rmse:0.516385+0.012395

```

```

print(tail(history))

```

```

##      train.rmse.mean train.rmse.std test.rmse.mean test.rmse.std
## 1:      0.482876      0.007863      0.507664      0.012382
## 2:      0.482116      0.007546      0.506943      0.012814
## 3:      0.481485      0.007599      0.506411      0.012642
## 4:      0.480852      0.007480      0.505880      0.012897
## 5:      0.480169      0.007492      0.505234      0.013035
## 6:      0.479464      0.007735      0.504545      0.012662

```

Validate Score

```

validate_predictions = predict(model, validate_matrix)
RMSE(validate_predictions, preparedValidateTarget)

```

```

## [1] 0.5618725

```

Train on full dataset

```

dtrain_final <- xgb.DMatrix(full_train_matrix, label = preparedFullTarget)

model_final = xgboost(dtrain_final, nround=100, nthread = 2, nfold = 5, metrics=list("rmse"),
                      max.depth = 3, eta = 0.1, objective = "reg:linear", print.every.n = 33)

## [0] train-rmse:3.926399
## [33] train-rmse:0.759091
## [66] train-rmse:0.609704
## [99] train-rmse:0.544402

history <- xgb.cv(data = dtrain_final, nround=100, nthread = 2, nfold = 5, metrics=list("rmse"),
                  max.depth = 3, eta = 0.1, objective = "reg:linear", print.every.n = 33)

## [0] train-rmse:3.926431+0.005475 test-rmse:3.926381+0.024544
## [33] train-rmse:0.755690+0.006431 test-rmse:0.762112+0.017506
## [66] train-rmse:0.608737+0.005373 test-rmse:0.618299+0.018470
## [99] train-rmse:0.546318+0.007077 test-rmse:0.558473+0.021436

```

Predict and un-log predictions

```

predictions = predict(model_final, test_matrix)
fixed_predictions = exp(predictions) - 1

```

Write the result

```

result = cbind(test2$datetime, fixed_predictions) %>% as.data.frame()
names(result) = c('datetime', 'count')
write.csv(result, 'submission.csv', quote = F, row.names = F)

```